

# Advanced Micro Devices

NASDAQ: AMD

A single-name forensic memo on narrative integrity, claim verification, insider distribution, and valuation.

Analysis window: April 22, 2026 to July 2, 2026, measured against trailing history

Engine: Walsh Narrative Decay Engine (MarketScholar AI)

Evidence base: daily forensic scorecard, SEC filing cross-reference, Form 4 insider record, 14 fiscal years of forensic accounting

Report date: July 2, 2026

**MP-2026-AMD-001 · CONFIDENTIAL · FOR INSTITUTIONAL DILIGENCE**

---

Prepared on the Walsh Narrative Decay Engine. Methodology published on SSRN (Forensic Narrative Classification and Equity Returns, abstract 6835001) and covered by USPTO provisional filings 63/971,470, 63/971,478, and 63/974,154. This document is research, not investment advice. See limitations and disclaimer on the final page.

## Executive summary

### CURRENT ENGINE VERDICT

## Narrative Risk

dominant window verdict: Narrative  
Trap (258 of 339 snapshots)

### CONFIDENCE

## 50

class: standard decay

### AS-OF PRICE

## \$518

July 2, 2026 · window high \$581 on  
June 30

AMD entered the window near \$303 on April 22, ran almost without interruption to an intraday high near \$581 on June 30 — a gain of more than 90 percent in ten weeks — and closed the window at about \$518. The central question for this memo is forensic, not directional: what does the engine read in that run, and does the conduct around it hold up. The answer has a different shape than a typical single-name review. The books are clean, the claims verify strongly, and the narrative is genuinely contested rather than an echo chamber. What stands out instead is the conjunction of a persistent elevated-suspicion regime, episodic coordination spikes, a valuation screen stretched to multiples of the model's estimate, and a documented \$164.5 million open-market insider sell program — including \$113.3 million by the Chair and CEO in two blocks placed squarely inside the run — against zero open-market insider purchases.

The current engine verdict is **Narrative Risk** at a confidence of 50, and the dominant daily verdict across the window was **Narrative Trap**. Unlike an exhausted name, AMD's narrative energy has not decayed: energy remaining held at ceiling throughout the window and the WKS score *rose* from roughly 20 to roughly 30, meaning the engine still sees fresh narrative force arriving. The base-case scenario assigns 47 percent to mean reversion within 30 days, 41 percent to continuation, and 12 percent to a volatility expansion — a nearly even split that is itself informative about how contested this tape is.

### Key findings, by confidence tier

**VALIDATED** Claim verification is strong and improved across the window: VMS 86–92, finishing at 91.8. The claims carrying the story tie to what the company files.

**VALIDATED** Forensic accounting is clean in the modern era: zero red flags in FY2025, Piotroski 0.875, conservative negative accruals. The 2012–2016 distress-era flags resolved a decade ago.

**VALIDATED** Insider distribution: \$164.5M in open-market sales across 67 transactions by five insiders, April 6 to June 15, with zero open-market purchases. CEO Lisa Su sold \$113.3M in two blocks, May 13 and June 10, both inside the run.

**VALIDATED** The narrative is contested, not crowded: 42 percent of 1,805 collision interactions are contradicting versus 33 percent amplifying — the opposite structure of an echo-chamber name.

**VALIDATED** A persistent elevated-suspicion regime: the suspicion score sat at 40–50 for most of the window rather than spiking episodically, with coordination pulses reaching 75 on April 30 and June 3.

**DIRECTIONAL** Roughly 3.7x the engine's modal fair value of about \$140. The model's band is severely unstable for AMD (artifact prints as low as \$8) and this is presented strictly as a screen, not a target.

**DIRECTIONAL** The insider sales carry no manipulation implication on their own. Rule 10b5-1 plan designations live in Form 4 footnotes the pipeline does not yet ingest; scheduled selling is the base case for officers at this level and is stated as such wherever the sales are discussed.

**What this memo does not claim.** It does not assert a price target and it does not make a short recommendation. It does not claim the insider sales were discretionary or informed; absent footnote-level plan data, it reports the calendar and the amounts, which are public record. It does not claim a per-name predictive track record for AMD, whose resolved-outcome sample inside the engine has not yet matured. These boundaries are detailed in Section 14.

## SECTION 02

# Methodology and data provenance

This report is produced by the Walsh Narrative Decay Engine, a forensic system that treats each market narrative as a measurable object with mass, energy, and a decay rate, and tracks how that object evolves as prices reveal whether the narrative was true. The underlying method is published on SSRN (Forensic Narrative Classification and Equity Returns, abstract 6835001), validated across a study of 187 equities over 17 months and 287,482 narrative observations, and is covered by three USPTO provisional filings.

For AMD the evidence base combines four independent layers: the daily forensic scorecard, SEC filing cross-reference, the Form 4 insider record, and fourteen fiscal years of forensic accounting history (FY2012–FY2025) — a history that, unusually, contains both a genuine distress era and the current hyperscale era, which sharpens the interpretation in Section 07.

| FORENSIC DIMENSION                  | DEPTH FOR AMD                            | WINDOW               | TIER        |
|-------------------------------------|------------------------------------------|----------------------|-------------|
| Daily forensic scorecard            | 52 daily snapshots (339 narrative rows)  | Apr 22 – Jul 2, 2026 | VALIDATED   |
| Coordination and collision          | Measured daily; 1,805 collision events   | Apr – Jul 2026       | VALIDATED   |
| Claim verification vs SEC           | Cross-referenced daily                   | Latest filings       | VALIDATED   |
| Insider transactions (Form 4)       | 78 non-derivative transactions           | Apr 1 – Jul 2, 2026  | VALIDATED   |
| Forensic accounting models          | 14 fiscal years                          | FY2012 – FY2025      | VALIDATED   |
| Keyword and collision drivers       | Measured in window                       | Apr – Jul 2026       | DIRECTIONAL |
| Relative value (fair value)         | Single-multiple framework, unstable band | Apr – Jul 2026       | DIRECTIONAL |
| 10b5-1 plan status of insider sales | Footnotes not ingested                   | n/a                  | EXCLUDED    |
| Per-name signal outcomes            | Not yet matured                          | n/a                  | EXCLUDED    |

## The two confidence tiers

Validated findings rest on direct measurement with adequate data: a value read off the filings, a Form 4 on the public record, or a daily-measured signal across the window. Directional findings are real and informative but depend on an assumption or a limited sample, and should not be read as precise. Section 14 also lists the metrics deliberately excluded from this report because they would mislead.

**Disclosed data-quality boundaries.** In the interest of diligence-grade honesty, the following were identified during preparation and handled as noted: (1) the daily scorecard series for AMD begins April 22, so no March baseline is claimed and all "window" statements refer to April 22 – July 2; (2) the fair-value framework is severely unstable for AMD, printing artifact values as low as \$8 on scattered days — sub-\$40

prints are treated as model failures and excluded, and the valuation section reports the modal estimate with the full band shown; (3) the energy-remaining metric held at its ceiling of 100 throughout the window, which the engine reads as continuous narrative refresh rather than a decay measurement — it is presented descriptively, not as evidence of narrative health; (4) the Beneish M-score is computable from AMD's XBRL extraction in only one fiscal year (FY2019, -2.13, clean), so no Beneish trend is claimed; (5) Form 4 footnotes, where Rule 10b5-1 plan designations appear, are not ingested — the plan status of every sale in Section 08 is therefore unknown to this system and stated as such; (6) per-name predictive outcomes have not matured and are not claimed.

# Price and narrative context

AMD entered the window at \$303.46 on April 22, rose almost monotonically through May — clearing \$400 by the first week of May and \$500 by late May — printed a window high of \$580.91 on June 30, and closed the window at \$517.82 on July 2. The move added roughly \$280 per share, better than 90 percent, in ten weeks, with only shallow pullbacks along the way.



Figure 1. AMD daily price across the analysis window (scorecard series). A near-uninterrupted advance from \$303 to a \$581 high, closing near \$518. The forensic interest is the conduct around the move, not the level.

|                                                                    |                                                                                                     |                                                                                                             |
|--------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| <div>AS-OF PRICE</div> <div>\$517.82</div> <div>July 2, 2026</div> | <div>MODAL MODEL FAIR VALUE</div> <div>\$140</div> <div>single-multiple screen; unstable band</div> | <div>INSIDER SALES IN WINDOW</div> <div>\$164.5M</div> <div>67 open-market sales · 0 open-market buys</div> |
|--------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|

For a forensic memo the level is less interesting than the manner of the move. The run occurred inside a persistent elevated-suspicion regime, with episodic coordination pulses; the narrative behind it kept refreshing rather than decaying; the coverage carrying it was contested rather than one-directional; and five insiders, led by the CEO, distributed into it on a documented calendar. The sections that follow take each of those threads in turn.

## SECTION 04

# Narrative integrity and coordination

AMD's integrity profile differs from the pattern the engine most often flags. Rather than a quiet baseline broken by a single suspicious episode, AMD ran a **persistent elevated-suspicion regime**: the suspicion score sat at 40 for most of the window, reached 50 in late April, and only decayed to 0-10 in the final week. Coordination was episodic on top of that floor, crossing the coordinated-pulse threshold with readings of **75 on April 30, 50 on May 8, 45 on May 26, and 75 on June 3**.




Figure 2. Coordination and suspicion scores, daily. The dashed line marks the coordinated-pulse threshold; the green band is the organic range. Red columns mark the CEO's two Form 4 sale blocks (May 13, June 10) for calendar reference only.

The forensically important reading is the regime, not any single day. A name that spends ten weeks with suspicion pinned at 40 while nearly doubling is a different object from a name that spikes once on a news event. At the same time, two honest qualifications cut the other way. First, the window's coordination pulses did not coincide tightly with the largest single-day price moves, so no pulse-to-price causation is claimed. Second, the readings at the close of the window are clean — coordination 0, suspicion 0-10, classified organic spread — so the regime is presented as a characteristic of the run, not a present-tense alarm.

**Scope note.** The scorecard series for AMD begins April 22, so no pre-run baseline is available and no claim is made about how the current regime compares with the stock's quiet state. The two red columns in Figure 2 mark the CEO's sale dates purely for calendar alignment with Section 08; they are not an assertion that the sales relate to the suspicion readings.

## SECTION 05

# Claim verification against SEC filings

This leg cuts in the company's favor, and it strengthened as the window progressed. The verification module extracts the claims embedded in the investment narrative and checks them against the company's filed financials, producing a verified-metric score (VMS). AMD's VMS ran between 86 and 92 across the window and finished at **91.8** — the hard numbers carried by the coverage tie to what AMD files.




Figure 3. Verified-metric (VMS), source-reliability (SRS), and narrative-risk (NRS) scores, daily. Claims verify at a consistently high level while narrative risk stays elevated — the caution in this memo is not about data integrity.

The claims the engine did flag were growth-narrative claims rather than filed-figure contradictions: statements that demand for the MI450 generation and Helios rack-scale systems is "exceeding early expectations," and that inferencing and agentic AI make the demand durable. These route to *needs-source* rather than *false* — they are forward-looking assertions that XBRL cannot adjudicate — and they are precisely the claims doing the narrative work at these prices. A reader performing diligence should treat the verified layer (revenue, margins, filed figures: strong) and the unverifiable layer (durability of AI demand: the entire bet) as separate objects, because the market is currently paying for the second.

**How to read this.** A high VMS means the coverage is not misquoting the filings. It does not mean the thesis is right. AMD's score says the inputs to the story are honest; whether the story's extrapolation holds is exactly what the valuation and insider sections are for.



## SECTION 06

# Valuation screen

The engine's modal fair value for AMD across the window is about **\$140 per share**, against an as-of price near \$518 — roughly **3.7 times** the model's estimate. That number is informative, but for AMD it must be read with even more care than usual, because the model's own behavior on this name is the least stable in this memo.




Figure 4. Price against the model fair value. The modal estimate is near \$140, but the framework printed artifact values as low as \$8 on scattered days. The band is shown in full; sub-\$40 prints are excluded from interpretation as model failures.

The fair value is produced by a single valuation framework, and across the window it printed a modal estimate near \$140 (with a secondary print at \$128) while also emitting scattered artifact values of \$8 to \$37 on days when a keyword-linked multiple misfired. We treat those prints as model failures and exclude them: a framework that can print \$8 for a company with AMD's cash flow is reporting its own breakage, not a valuation. What survives the cleanup is still a stark statement — the market is paying well over three times what the framework's multiple implies — but it is a **one-variable finding**, and AMD's premium plausibly reflects factors a single multiple cannot carry: the MI450/Helios product cycle, hyperscaler capex commitments, and the optionality the market assigns to AI compute share gains against the incumbent.

**How to use this.** Read Section 06 as: on the engine's framework, AMD screens as one of the most richly valued names in coverage, by a wide and persistent margin, and the model's instability on this name is itself disclosed. Do not read \$140 as a price objective or a basis for a short. The defensible restatement is that the current price embeds an AI-demand durability assumption that no filed figure can verify — which is what connects this section to Section 05.

## Forensic accounting corroboration

The fundamental forensic models look for the fingerprints of earnings manipulation and financial distress. AMD offers fourteen fiscal years of history, and that history has a shape few large caps can show: a genuine distress era, fully resolved, followed by the current era of clean prints. The table shows the arc.

| FACTOR                           | FY2015<br>(DISTRESS ERA) | FY2022 | FY2024 | FY2025 | READING                                   |
|----------------------------------|--------------------------|--------|--------|--------|-------------------------------------------|
| Red-flag count                   | 5                        | 1      | 1      | 0      | Distress-era flags fully resolved         |
| Accruals ratio                   | −0.14                    | −0.03  | −0.02  | −0.04  | Negative; conservative earnings quality   |
| Piotroski F-score                | 0.17                     | 0.50   | 1.00   | 0.875  | From distressed to top-decile quality     |
| Negative equity                  | Yes                      | No     | No     | No     | 2015 episode only                         |
| Cash burn                        | Yes                      | No     | No     | No     | Burned cash FY2012–FY2015; positive since |
| Goodwill + intangibles, % assets | 9%                       | 71%    | 63%    | 54%    | Xilinx acquisition; amortizing down       |

Three observations matter. First, the modern record is clean: FY2025 carries zero red flags, accruals are consistently negative (cash flow runs ahead of reported earnings), and Piotroski sits at the top of its range. Second, the balance sheet still carries the Xilinx acquisition — goodwill and intangibles peaked at 71 percent of assets in FY2022 and have amortized down to 54 percent, which is not a red flag but concentrates carrying value in assets whose worth depends on continued performance. Third, the distress era is relevant context for the insider section: the same franchise that had negative equity in 2015 now supports nine-figure insider liquidity events, which is a legitimate wealth-realization motive that cuts *against* reading the sales darkly.

**Disclosure on Beneish.** The Beneish M-score is computable from AMD's XBRL extraction in only one fiscal year (FY2019: −2.13, no manipulation signal). The inputs required for the eight-ratio model do not resolve cleanly from AMD's tagging in other years, so no Beneish trend is claimed for this name. The manipulation assessment rests on the accruals, Piotroski, and flag-count evidence above, all of which point the same clean direction.

## SECTION 08

# Insider transactions: the distribution record

This is the section that most differentiates AMD's window, and it is built entirely from Form 4 filings on the public record. Between April 6 and June 15, 2026, five AMD insiders executed **67 open-market sales (Code S, non-derivative) totaling \$164.5 million**. Over the same period the record shows **zero open-market purchases** — the only acquisitions are option exercises (Code M) and two gift transactions.

| INSIDER             | TITLE                     | SALE DATES               | TXNS | TOTAL    |
|---------------------|---------------------------|--------------------------|------|----------|
| Lisa T. Su          | Chair, President & CEO    | May 13 · Jun 10          | 49   | \$113.3M |
| Mark D. Papermaster | CTO & EVP                 | Apr 6 – Jun 15 (6 dates) | 6    | \$26.5M  |
| Paul D. Grasby      | EVP & Chief Sales Officer | May 8                    | 1    | \$10.8M  |
| Forrest E. Norrod   | EVP & GM, Data Center     | May 20                   | 9    | \$8.4M   |
| Nora Denzel         | Director                  | May 29 · Jun 2           | 2    | \$5.5M   |



Figure 5. Open-market insider sales (bars, right axis) against the share price (grey line, left axis). The CEO's two blocks — \$55.7M on May 13 and \$57.6M on June 10 — land inside the steepest phase of the advance.

The calendar is the finding. The CEO's first block, \$55.7 million, printed on May 13 with the stock around \$410, mid-run; the second, \$57.6 million, printed on June 10 around \$510, three weeks before the \$581 high. The CTO sold on six separate dates spanning the entire window. The Chief Sales Officer and the head of the Data Center business — the division carrying the AI narrative — each sold eight-figure blocks in May. Whatever else is true, the officers closest to the MI450 and Helios demand story monetized into the strongest tape their company has ever printed, and no insider bought a share on the open market against them.

**The mandatory caveat, stated prominently.** Rule 10b5-1 plan designations appear in Form 4 footnotes that this system does not yet ingest. Officers at this level overwhelmingly sell through pre-set plans, and a plan adopted before the run would make this scheduled selling — diversification on a calendar, not a view on the stock. This memo therefore claims the amounts, the dates, and the zero-buy asymmetry, all of which are public record, and claims nothing about intent. A diligence team should pull the footnotes on the May 13 and June 10 filings before drawing any adverse inference; if the plans predate the run, this section reads as routine wealth management at unusual scale, and that reading should be preferred.

## SECTION 09

# Narrative dynamics over time

The Walsh engine models each narrative as a physical object with mass and energy that decays as price reveals whether the story held. AMD's panels tell a story of refresh, not decay — the structural opposite of an exhausted name.




Figure 6. Narrative mass, energy remaining, and WKS score, daily. Energy held at ceiling throughout (see disclosure); WKS rose across the window rather than collapsing.

Narrative mass held in a healthy band and the WKS score, far from collapsing, climbed from roughly 20 at the window open to roughly 30 at the close. In the engine's vocabulary, fresh narrative force kept arriving faster than price could exhaust it: each leg of the run was met by a new claim layer — MI450 demand, Helios systems, hyperscaler capex, AI power constraints — that reset the field. This is why the verdict is Narrative *Risk* rather than Exhausted Narrative: the engine is not saying the story has run out; it is saying the story is doing an unusually large share of the valuation work while insiders distribute into it.

Narrative tone by day (green bullish, red bearish, grey mixed)




Figure 7. Narrative tone by day. Bullish tone dominates the run, with mixed readings clustering at the window close.

**Reading these series.** The energy-remaining metric held at its ceiling value of 100 for the full window. The engine interprets a pinned ceiling as continuous narrative refresh, but a pinned metric carries less information than a moving one, so it is shown descriptively and no energy-based claim appears in the findings tiers. WKS and narrative mass are proprietary Walsh-engine measures, presented as observed trajectories to characterize the narrative field rather than as standalone trade signals.

## SECTION 10

# Coverage and source landscape

Across the window the engine attributed AMD coverage to the publications below. The table shows article volume and the bullish-versus-bearish split for the most active outlets.

| PUBLICATION        | ARTICLES | BULL | BEAR |
|--------------------|----------|------|------|
| The Motley Fool    | 156      | 76   | 36   |
| Benzinga           | 67       | 22   | 13   |
| Barron's           | 42       | 15   | 17   |
| CNBC               | 32       | 15   | 0    |
| AOL.com            | 24       | 10   | 0    |
| The Globe and Mail | 16       | 2    | 2    |
| Seeking Alpha      | 16       | 6    | 2    |
| Yahoo Finance      | 15       | 9    | 2    |
| MarketWatch        | 14       | 6    | 1    |
| TIKR.com           | 13       | 12   | 0    |

Two observations matter. First, the aggregate tilt is bullish — the two largest rated outlets run roughly two-to-one bullish, and several mid-size outlets (CNBC, AOL, TIKR) published essentially no bearish pieces at all during a 90 percent advance. Second, the tilt is not uniform: Barron's ran net-bearish, and Motley Fool's 36 bearish pieces are a real dissenting mass. This mixed-but-tilted structure matches the collision composition in Section 12 — AMD's coverage is a genuine argument that the bulls are currently winning, not a manufactured chorus. In the engine's experience the argument structure is healthier than an echo chamber, but it also means the marginal buyer is being supplied with a steady stream of affirmation at the highs.

**On named-source accuracy.** The engine computes per-author hit rates, but for most authors the sample is a handful of articles, far too small to support a reliability judgment about any individual. Coverage volume and tone are reported as factual; per-author accuracy is treated as descriptive only and no named person is evaluated.

## SECTION 11

# Narrative drivers: keywords

The keyword layer measures the average forward price move after a given word appears in an AMD headline. AMD's information environment is saturated with artificial intelligence: **ai\_mention appears in 962 headlines**, by far the most frequent driver, and the AI-adjacent themes dominate the high-impact tail.




Figure 8. Average 5-day price move after each keyword appears in a headline, with article counts. Green positive, red negative. Several keywords rest on small samples.

The pattern is a melt-up tape. Nearly everything is followed by gains: earnings beats (+4.6 percent), price-target changes (+7.8), upgrades (+7.0), AI power demand (+13.0 on 33 articles) — and, tellingly, even **analyst downgrades were followed by +14.1 percent** and "AI capex bubble" mentions by +11.4 percent. When bearish framings are systematically followed by rallies, the tape is telling you that skepticism itself has become a contrarian entry signal in the current regime — which is the behavioral signature of a momentum phase, not of a market weighing evidence. The lone meaningful negative is partnership\_deal at -1.9 percent, a mild sell-the-news effect on the window's many ecosystem announcements.

**How to weigh this.** Keyword price-impact values are real but several rest on small article counts, which makes any single figure noisy and raises a multiple-comparisons risk across the full set. The +14.1 percent after downgrades (n=27) should be read as texture about regime behavior, not as a tradeable edge. One keyword (GLP1\_disruption, n=12) is a sector-tagging artifact unrelated to AMD and is excluded from the chart.



## SECTION 12

# Collision forensics and credibility

Collision forensics examine how individual narrative events interact — whether they reinforce one another, contradict one another, or arrive in coordinated clusters. For AMD the picture is contested, and that is the single most important structural difference between this name and an echo-chamber narrative.




Figure 9. Composition of 1,805 collision interactions in the window. A contradiction-heavy structure is the fingerprint of a genuine two-sided debate.

Of 1,805 collision interactions, **761, or 42 percent, are contradicting** — narratives pushing against each other — versus 599, or 33 percent, amplifying. Coordinated clusters account for 126 events, or 7 percent, concentrated around the April 30 and June 3 pulses documented in Section 04, and high-volume single-source events (one outlet flooding the zone) account for a further 15 percent. A contradiction-dominant structure means the AI-demand thesis is being actively argued, not merely repeated: bears on valuation, memory-supply constraints, and capex sustainability collide daily with bulls on MI450 demand and share gains. In the engine's framework this is healthier than an amplifying chorus — but it also means the current price is the market's live verdict on a contested argument, with the contest still open. The coordinated-cluster share of 7 percent is elevated enough to note and consistent with the episodic pulses, but it is not the dominant mode.

## SECTION 13

# Options positioning and market structure

The final descriptive layer covers options positioning and dark-pool flow, both carrying distinctive readings at the window close.

### PUT/CALL SKEW

**1.13**

put OI 1.57M vs call OI 1.39M —  
hedged, not euphoric

### MAX PAIN

**\$325**

vs \$518 price — positioning far  
below spot

### DARK POOL (LATEST SAMPLE)

**Buying**

36.2% of volume

Two of the three readings cut against the simplest bearish framing, and they are reported accordingly. Options positioning is put-heavy (skew 1.13), meaning the market is paying for downside protection rather than chasing calls — at a top built on euphoria one usually finds the opposite, and heavy put walls can cushion declines. The latest dark-pool sample reads as buying at 36 percent of volume, which leans constructive. On the cautionary side, max pain at \$325 sits nearly \$200 below spot, an extreme gap: to the extent expiry mechanics exert any pull, it is downward, and the gap itself is a measure of how far price ran ahead of where the options market accumulated its positioning. Earnings arrive in roughly 33 days (early August), which is the natural resolution point for the contested narrative documented throughout.

**Data coverage.** Options and dark-pool readings are end-of-window samples, not full-window series, and characterize positioning as of July 2 only. The dark-pool direction flips session to session and is a secondary input, not a load-bearing finding.

## SECTION 14

# Signal summary and confidence appendix

The engine's daily verdict for AMD was cautionary on nearly every session of the window: **Narrative Trap on 258 of 339 snapshots**, Narrative Risk on 69, with scattered Drift Compression Short, Monitoring, and Omission Cascade states. There was no Accumulate day. The current verdict is Narrative Risk at confidence 50. The scenario distribution as of the report date assigns 47 percent to mean reversion within 30 days, 41 percent to continuation, and 12 percent to a volatility expansion.

## Engine-wide validated signals (not name-specific)

The performance evidence below is drawn from the engine's full universe and its published backtest registry. It is cited to characterize the system, not AMD. No per-name predictive accuracy is claimed for AMD, whose resolved-outcome sample has not yet matured.

| SIGNAL             | PERFORMANCE     | SAMPLE  | SIGNIFICANCE |
|--------------------|-----------------|---------|--------------|
| EXHAUST_HIGHCONF   | Sharpe 2.25     | n = 459 | p = 0.0003   |
| FWOR_VMS           | Sharpe 2.37     | n = 117 | significant  |
| DUAL_CONFLUENCE_T1 | +3.65% 5d alpha | n = 417 | significant  |

## Every material claim, mapped

| CLAIM IN THIS MEMO                   | TIER        | BASIS / CAVEAT                                     |
|--------------------------------------|-------------|----------------------------------------------------|
| Claims verify against filings        | VALIDATED   | VMS 86-92 across window, 91.8 at close             |
| Modern accounting record clean       | VALIDATED   | 0 flags FY2025; accruals -0.04; Piotroski 0.875    |
| \$164.5M insider sales, zero buys    | VALIDATED   | 67 Code-S Form 4 transactions, Apr 6 - Jun 15      |
| CEO sold \$113.3M inside the run     | VALIDATED   | Blocks May 13 (\$55.7M) and Jun 10 (\$57.6M)       |
| Persistent elevated-suspicion regime | VALIDATED   | Suspicion 40-50 most of window; pulses to 75       |
| Narrative contested, not crowded     | VALIDATED   | 42% contradicting vs 33% amplifying, n = 1,805     |
| Coverage tilt bullish but mixed      | VALIDATED   | Fool 76/36, CNBC 15/0, Barron's net-bearish        |
| ~3.7x modal fair value               | DIRECTIONAL | Single multiple; band unstable, artifacts excluded |
| Melt-up keyword regime               | DIRECTIONAL | Downgrades +14.1% (n=27); small samples            |
| Options/dark-pool readings           | DIRECTIONAL | End-of-window samples only                         |
| Insider sales were discretionary     | EXCLUDED    | 10b5-1 footnotes not ingested; no intent claim     |
| Beneish M-score trend                | EXCLUDED    | Computable in one year only (FY2019, clean)        |
| Energy-based narrative health        | EXCLUDED    | Metric pinned at ceiling; low information          |

|                                     |          |                            |
|-------------------------------------|----------|----------------------------|
| Pre-window coordination baseline    | EXCLUDED | Series begins Apr 22       |
| Fair value as target or short basis | EXCLUDED | Screen only                |
| Per-name predictive accuracy        | EXCLUDED | Outcome sample not matured |

Data as of July 2, 2026. Daily forensic scorecard across the April 22 – July 2, 2026 window (52 daily snapshots, 339 narrative rows). Forensic accounting: fourteen fiscal years, FY2012–FY2025. Insider record: Form 4 filings, April 1 – July 2, 2026.

## SECTION 15

# Limitations and disclaimer

**Not investment advice.** This report is forensic research intended for institutional diligence. It does not constitute investment advice, a recommendation, or an offer or solicitation to buy or sell any security. Nothing here is a price target, a short recommendation, or a forecast of returns.

**On the insider record.** All insider transaction data is drawn from Form 4 filings on the public record. This memo reports amounts, dates, and the absence of open-market purchases. It makes no claim that any sale was discretionary, informed, or improper, and it discloses that Rule 10b5-1 plan designations — which would establish the sales as pre-scheduled — reside in filing footnotes this system does not ingest. Readers must verify plan status from the filings before drawing any inference about intent.

**Methodology limitations.** The Walsh Narrative Decay Engine measures narrative structure, claim consistency, and relative value. It does not model every fundamental, macro, or idiosyncratic driver of price. The fair-value framework in Section 06 is a single-multiple screen, is explicitly directional, and was materially unstable on this name; artifact prints were excluded and disclosed. Per-name predictive outcomes have not matured, and all performance statistics cited are engine-wide.

**Data and sourcing.** Figures are drawn from the engine's pipeline as of July 2, 2026, including SEC filing data, Form 4 records, market price data, and the narrative observation base. Six data-quality boundaries are disclosed in Section 02 and reflected in the confidence tiers throughout. Readers should weigh validated and directional findings differently.

**Confidentiality.** This document is confidential and prepared for the named recipient. Methodology is published on SSRN (abstract 6835001) and covered by USPTO provisional filings 63/971,470, 63/971,478, and 63/974,154.

MarketScholar AI · Walsh Narrative Decay Engine · Case MP-2026-AMD-001 · July 2, 2026